

# Software Defined Networking and OpenFlow


Marco Ruffini: [marco.ruffini@tcd.ie](mailto:marco.ruffini@tcd.ie)

# Lecture content

- SDN definitions
- OpenFlow table: matching, actions and monitoring
- Control plane command line
- Managing QoS through meters and queues

# SDN: Definitions, Concepts, and Terminology

- › **Data plane:** network infrastructure consisting of interconnected forwarding devices (a.k.a., forwarding plane).
- › **Forwarding devices:** data plane hardware- or software devices responsible for data forwarding.
- › **Flow:** sequence of packets logically belonging together (e.g. based on source-destination pair); flow packets receive identical service at forwarding devices.
- › **Flow rules:** instruction set that act on incoming packets (e.g., drop, forward to controller, etc)
- › **Flow table:** resides on switches and contains rules to handle flow packets.




Source:

"Software-Defined Networking: A Comprehensive Survey",  
Kreutz et al., <https://arxiv.org/pdf/1406.0440>.

# SDN: Definitions, Concepts, and Terminology

- › **Southbound interface:** (instruction set to program the data plane) + (protocol between control- and data planes).
- › **Control plane:** controls the data plane; logically centralized in the “controller” (a.k.a., network operating system).
- › **Northbound interface:** API offered by control plane to develop network control- and management applications.
- › **Management plane:** functions, e.g., routing, traffic engineering, that use control plane functions and API to manage and control network infrastructure



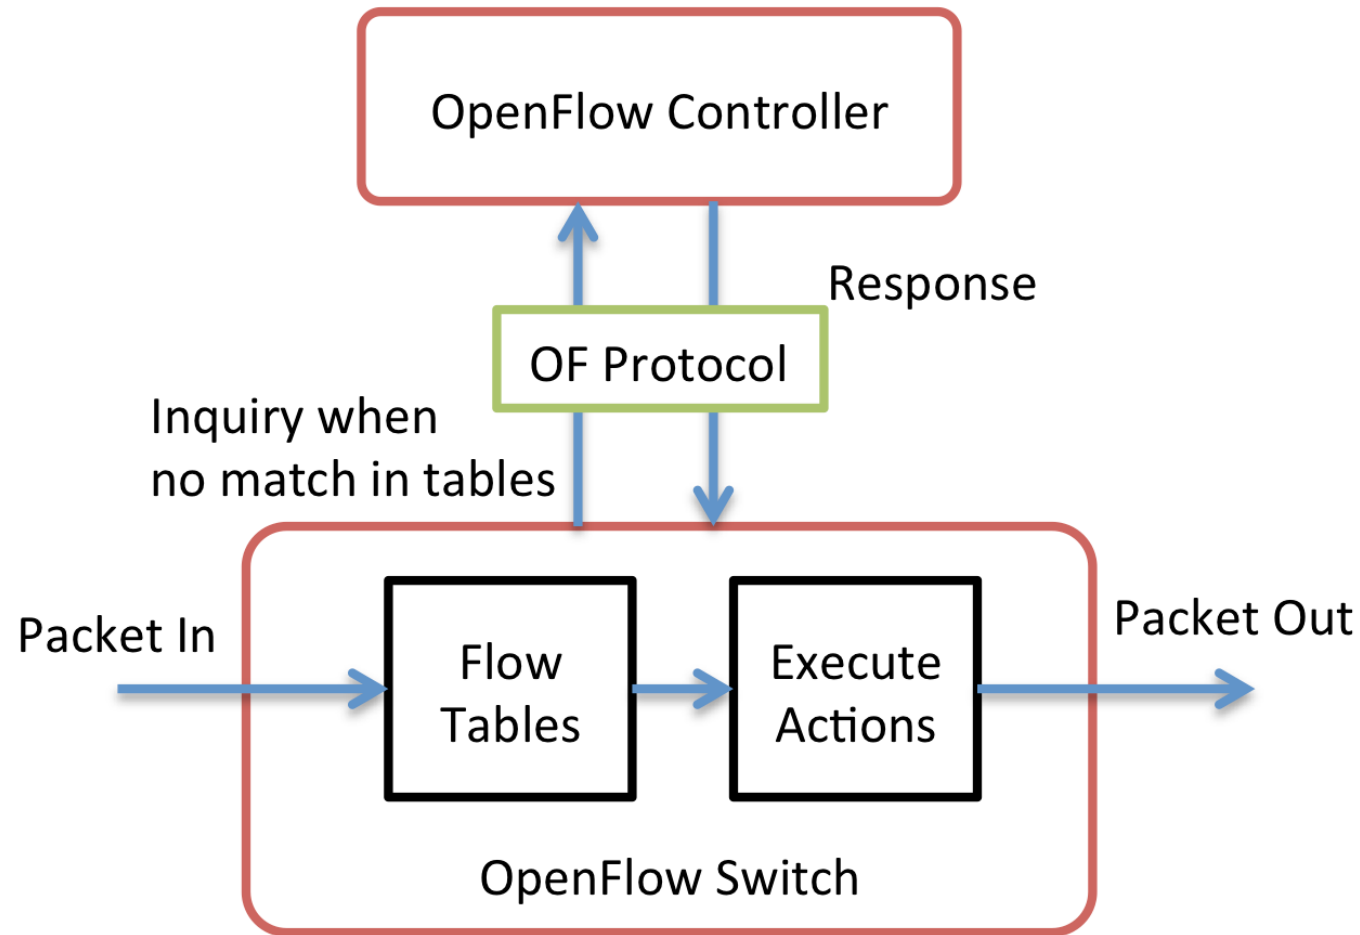
Source:

“Software-Defined Networking: A Comprehensive Survey”,  
Kreutz et al., <https://arxiv.org/pdf/1406.0440>.

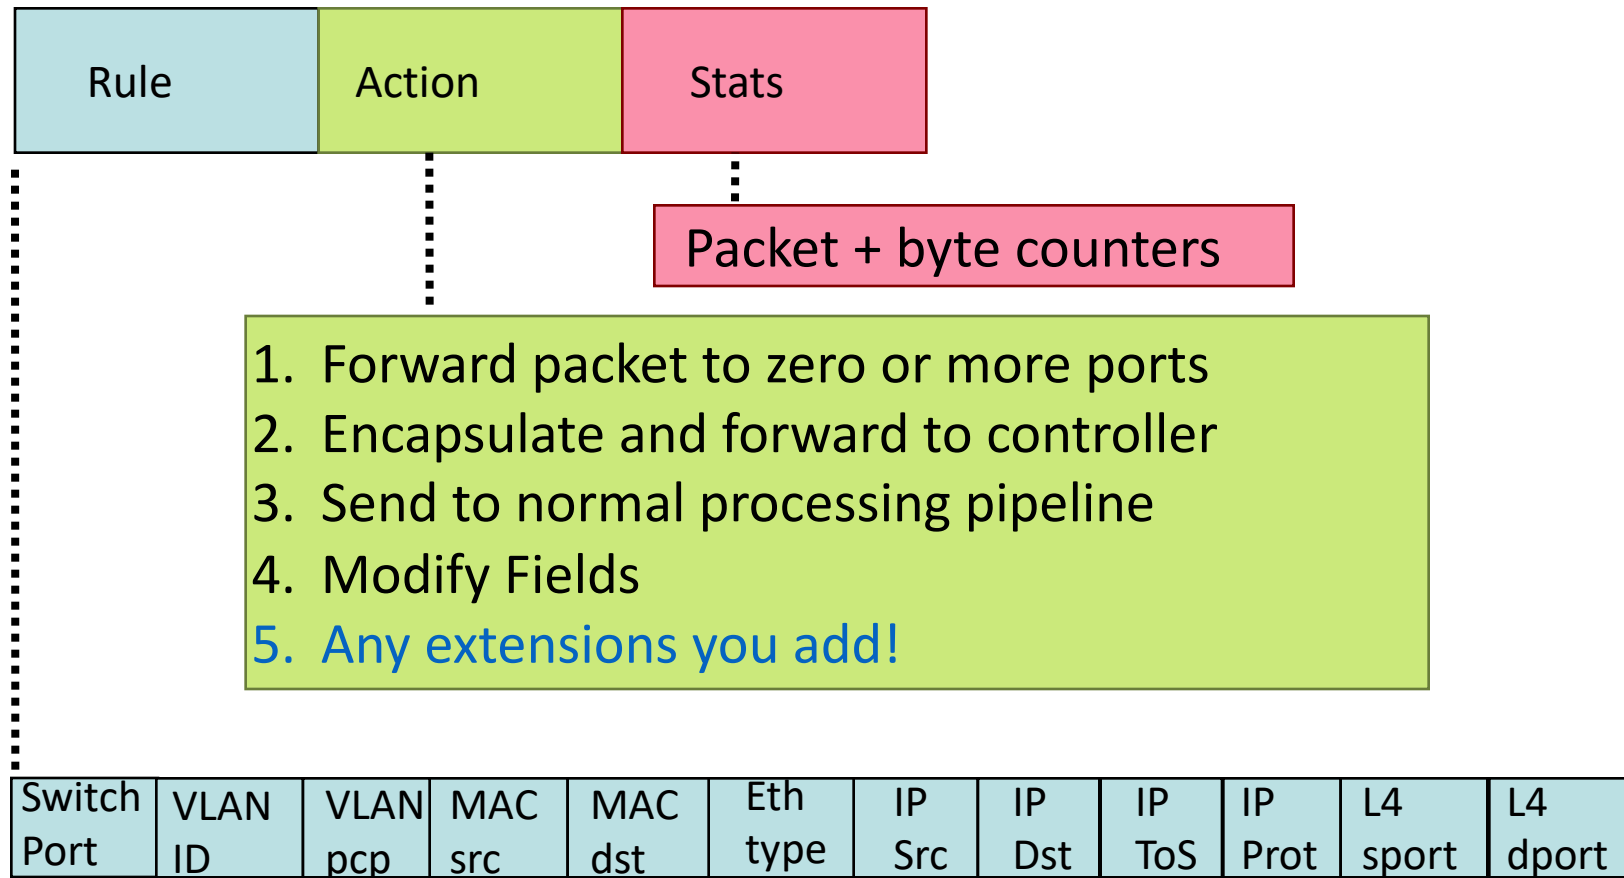
# Timescales

	Data	Control	Management
Time-scale	Packet (nsec)	Event (10 msec to sec)	Human (min to hours)
Tasks	Forwarding, buffering, filtering, scheduling	Routing, circuit set-up	Analysis, configuration
Location	Line-card hardware	Router software	Humans or scripts

# OpenFlow Flow Table Entries



# OpenFlow Flow Table Entries



+ mask what fields to match

# Examples

## Switching

[illegible]

## Flow Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
port3	00:20:..	00:1f:..	0800	vlan1	1.2.3.4	5.6.7.8	4	17264	80	port6

## Firewall

[illegible]



# Examples

## Routing

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	5.6.7.8	*	*	*	port6

## VLAN Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f..	*	vlan1	*	*	*	*	*	port6, port7, port9

# Flow Routing vs. Aggregation

Both models are possible with OpenFlow

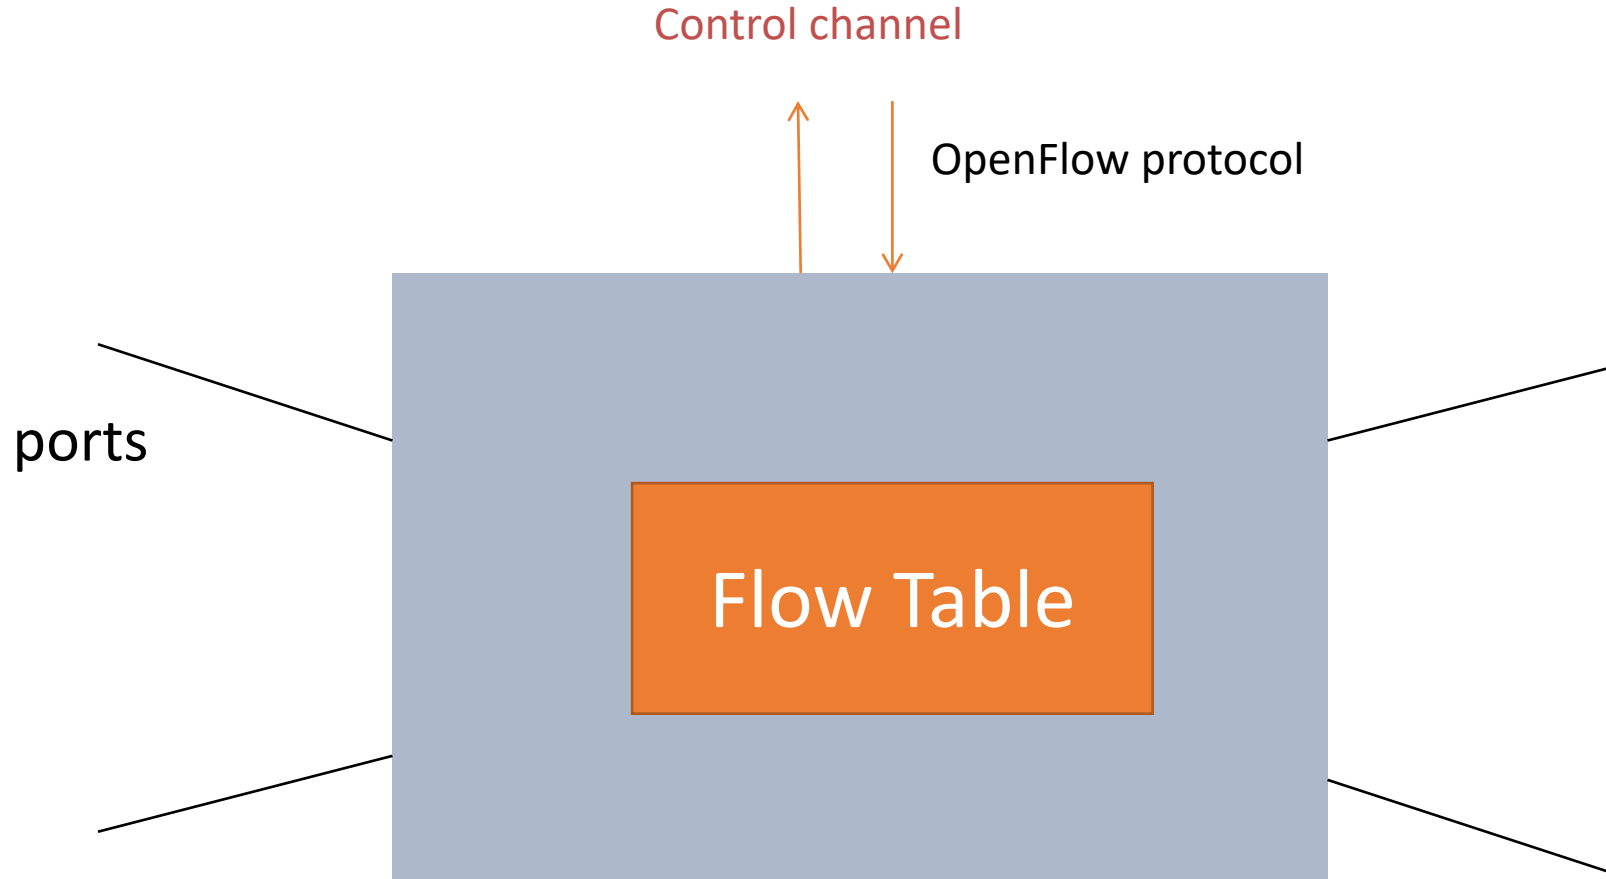
## Flow-Based

- Every flow is individually set up by controller
- Exact-match flow entries
- Flow table contains one entry per flow
- Good for fine grain control, e.g. campus networks

## Aggregated

- One flow entry covers large groups of flows
- Wildcard flow entries
- Flow table contains one entry per category of flows
- Good for large number of flows, e.g. backbone

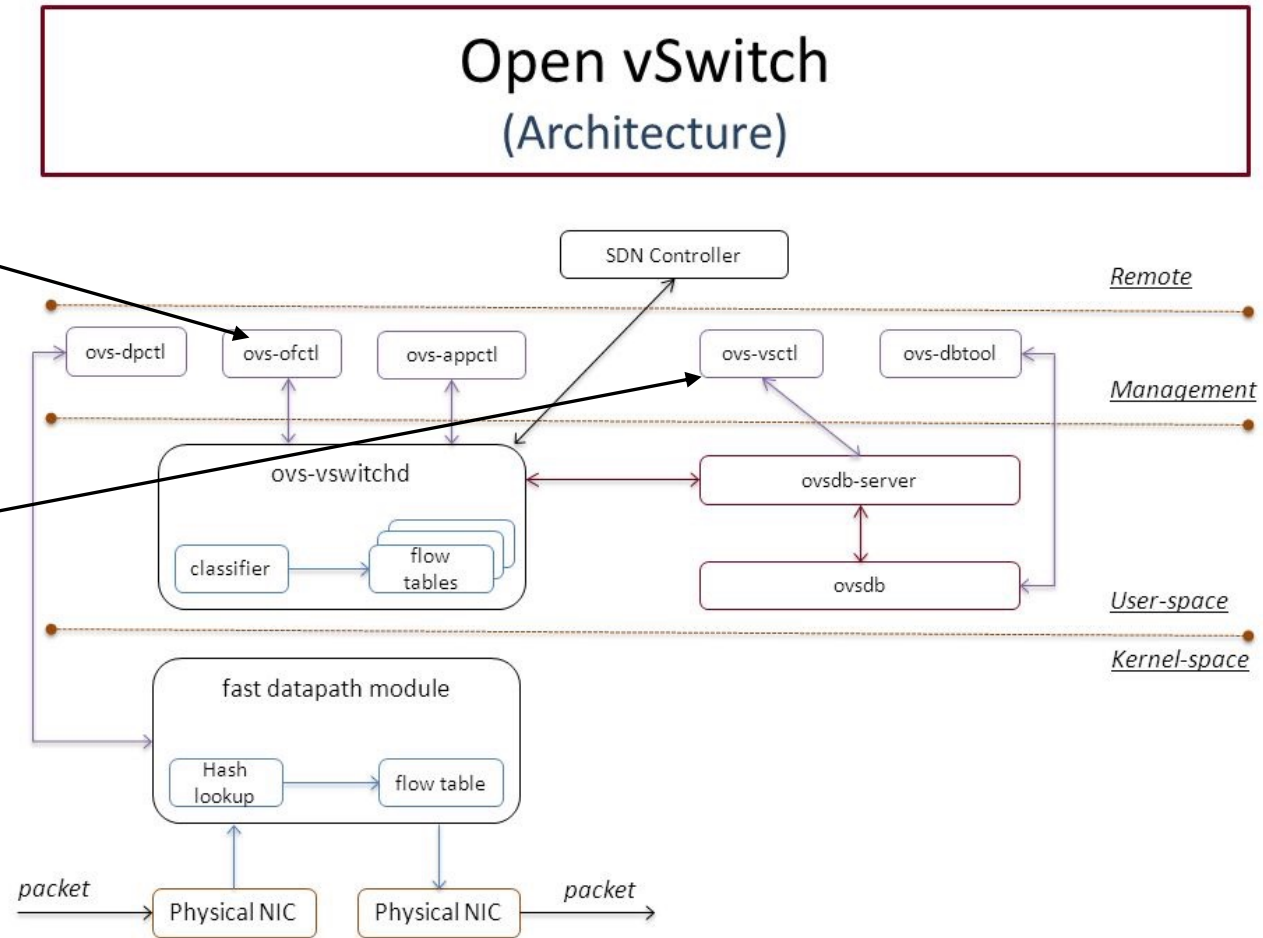
# OpenFlow Switch



# OvS architecture

We will use ovs-ofctl commands to control flow tables

We will use ovs-vsctl for virtual switch management



# Control plane commands

- The controller can provide a range of control messages to the switch.
- We discuss these APIs with reference to the Open vSwitch(OvS) (<https://www.openvswitch.org/>)
  - ofctl: command set for controlling openflow switches → commands operating on flow tables (flow rules, monitoring, etc)
  - vsctl: command set for administering the OvS → commands operating on switch management/configuration (i.e., ports, queues,...)

# Flow table command line (ovs-ofctl)

- **ovs-ofctl dump-flows [switch]:** shows all flow entry on the device called [switch] → this lets you know all rules currently installed in the switch
- **ovs-ofctl del-flows [switch] [flow]:** deletes the flow entry on the device called [switch] that matches [flow]; if the argument [flow] is omitted, it deletes all flows in the switch → after this command the flow table will be empty and the switch will have no information on how to switch incoming packets
- **sudo ovs-ofctl add-flow [switch] [match rules] [action rules]:** installs the flow rule on [switch]; any flow whose header will match the [match rules] will be handled according to the [action rules] example: *[match rules]=dl src=00:00:00:00:00:02,dl type=0x806; [action rules]: actions=output:"s1-eth1"*
- **sudo ovs-ofctl mod-flow [switch] [match rules] [action rules]:** modify the match and action rules of an existing flow.

# Other important match and action rules (ovs-ofctl)

- Matching:

- *in\_port=port* (true if packet comes in from port)
- *dl\_vlan=vlan* (true if header VLAN matches vlan)
- *dl\_src=xx:xx:xx:xx:xx:xx* (true if source MAC address matches xx:xx:xx:xx:xx:xx )
- *dl\_dst=xx:xx:xx:xx:xx:xx* (true if dest MAC address matches xx:xx:xx:xx:xx:xx )
- *dl\_type=ethertype* (true if ETHERTYPE bits match ethertype). – <https://en.wikipedia.org/wiki/EtherType>
- *nw\_src=ip[/netmask]* (when Ethertype is set to 0x800 this matches on source IPv4 address and mask)
- *nw\_dst=ip[/netmask]* (when Ethertype is set to 0x800 this matches on destination IPv4 address and mask)
- *tcp\_src=port; tcp\_dst=port; udp\_src=port; udp\_dst=port* (match on TCP or UDP ports)

- Actions:


- *output:port* (sends the packet out on port)
- *enqueue:port:queue\_ie* (send the packet to a specific QoS queue for that port – the queue needs to be configured separately in ovs-vsctl)
- *all* (sends a copy of the packet on all ports, except where it came from)
- *in\_port* (sends the packet out on the port where it came in)
- *controller* (sends the packet to the controller)
- *drop* (drops the packet) – this is the default rule if no matching rule is installed
- *meter:meter\_id* (apply the meter rule written in rule\_id to this packet)
- many more commands, i.e., modify fields, add tags, etc...

See <https://www.openvswitch.org/support/dist-docs-2.5/ovs-ofctl.8.txt>

# Managing QoS with meters

OpenFlow 1.3 introduces meters:

- A meter provides a threshold, after which some action is carried out.
- Meters can be applied as part of an action on a flow entry. They are based on a given threshold expressed in kbps or packets per second.
- The OvS implementation is quite simple, so only one threshold can be set in the meter and any packet above that level is dropped
- Remember that a meter works by linking it to a specific flow rule. Multiple flow rules can link to the same meter.
- **It basically implements a policer\***




\*Meters however also allow to set a maximum burst rate size



# Managing QoS with queues

- Queues are not specific to OpenFlow but part of OvS itself.
- They are managed through the ovs-vsctl commands rather than the OpenFlow ones (ovs-ofctl)
  - Multiple queues can be linked to the same output port with different min and max rate
  - Flow rules can then be sent to a specific queue ID on a given port
- This act more like as scheduler, although the only configuration we can provide is committed (min-rate) and peak (max-rate)



# Managing QoS through meters (ovs-ofctl)

*Meters can be used from ovs-ofctl to set a maximum available rate for packets matching a specific rule.*

- *add-meter [switch] [meter]: adds a new meter in the switch with details expressed in [meter]*
- *mod-meter [switch] [meter]: modifies an existing meter*
- *del-meter [switch] [meter]: deletes an existing meter*
- *dump-meter [switch] [meter]: shows a list of existing meters*
- *meter-stats [switch] [meter]: reports statistical information on a given meter*
- *[meter] fields:*
  - *meter=id (gives an identifier to this meter)*
  - *Kbps or pktps (whether it operates on kbps or packets per second)*
  - *stats (if we want to collect packet statistics)*
  - *burst (is there is a burst size for all bands)*
  - *bands=band parameters*
    - *type=type (only drop available)*
    - *rate=value (the rate of this meter)*
    - *burst\_size=size (the max burst size, if any for this specific band)*

# Managing QoS through queues (ovs-vsctl)

- ovs-vsctl allows to define QoS also for specific queue, in the port management functionality.
- **ovs-vsctl set port [port] [configs]:** *(set relevant parameters for port configuration)*
- Typical example for setting up queue management:
  - **Qos config parameter:** `qos=@newqos -- --id=@newqos create qos type=linux-htb queues=0=@q0 -- --id=@q0 create queue other-config:min-rate=500000 other-config:max-rate=1000000`
  - **Example with two queues:** `$ sudo ovs-vsctl set port s1-eth3 qos=@newqos -- --id=@newqos create qos type=linux-htb queues=0=@q0,1=@q1 -- --id=@q0 create queue other-config:min-rate=200000000 other-config:max-rate=500000000 -- --id=@q1 create queue other-config:min-rate=50000000 other-config:max-rate=100000000`
  - Key arguments are:
    - *min-rate: this is the assured rate that the queue needs to provide*
    - *max-rate: this is the maximum rate for the queue (packets above this rate are dropped)*
- many more commands: <https://www.openvswitch.org/support/dist-docs/ovs-vsctl.8.txt>